

Predicting coding question quality using Stack Overflow ratings

Slide 2: Overview

- Stack Overflow (SO)
- What is a question?
- Support Vector Machine (SVM)
- Methodology
- Experiments and Results
- Summary
- Demo

Slide 3: Stack Overflow

Stack Overflow (SO) was created by Jeff Atwood and Joel Spolsky, and was released in September 2008. They created it to offer programmers a place where they could ask questions and get answers to the questions. To be able to measure quality, they used votes, where users could up-vote questions they found useful, and down-vote questions that were wrong. Furthermore, a question is not restricted to have only one answer, allowing multiple solutions to be present.

In fact, Stack Exchange (SE), which were released one year later, is built upon the same model. This means that if you develop a system for SO, you could also expand it to cover all of SE.

SO uses gamification to reward the users for their participation. Gamification means that you use game elements in places which normally would not be considered a game. In SO, users can be rewarded in various ways, but the three main elements are votes, reputation and badges.

Votes are used as a measurement of the question (and answers) quality and usefulness, but are only shown on the given post. SO also sorts the questions and answers based on score, with the exception being answers that are marked as accepted. An accepted answer are answers which the questioner found to be correct for their problem.

Reputation and Badges follows the users, where one can compare Badges to achievements in games. Reputation is not only used to show how much you have participated, but it also restrains the amount of freedom you have on the site (e.g. commenting, voting, answering, etc). Reputation can in fact be used as a measurement of expertise, because there is a limit to how much reputation you can earn daily. To draw an example from games, it is not like in World of Warcraft or Tera, where you can grind repetitive quests somewhat endlessly.

Users can only earn up to 200 points of reputation each day, where an answer gives +10, and question +5. That amounts to posting 20 answers, or 40 questions. After reaching the daily cap, the only way users can earn more reputation, is by having their answer marked as accepted or earning bounties. Bounties are a currency system. If a user has a question no one answers, or the answers given does not solve their problem, users can trade parts of their reputation for a solution.

I could not find a good identifier for what a good question is. There is what I would call a bias factor. By this I mean that if a certain amount of people share the same problem, it becomes a good

question. Not because of the question asked, but because of the problem it needs to solve. However, these are in most cases weighed up by using the Wiki feature, namely Community Wiki. Community Wiki are usually added to posts that are considered to be helpful to the community (and you can even search for Wiki posts in the search field).

The bad questions were easier to spot. A common denominator was the lack of effort when asking their questions. For the most code related questions, they added a code snippet, and said "This doesn't work. Why?". Other examples were large code examples showing a lot of code that was not related, or no code at all and just showing the error message.

If you went to your teacher with 50 lines of code, and said "This doesn't work. Why?". Do you think s/he would know what was wrong? In most cases, the obvious reply would be "what have you tried?", "what is the expected output" or "what is the error that you are getting". A lot of code examples were also badly formatted and had syntax error. There were also obvious signs of homework/school related topics, and in one question the code even contained the namespace "assignment".

SO also wants questions to be unique, so duplicates are often down-voted. There are however some cases where this does not happen, and it could therefore be interesting to see in the long term which would be the most viewed.

TODO: Look into *"... how much it would have cost to pay people to rate the questions if you were doing it on a time spend reading and rating"*

Slide 4: Picture of good SO question

"How do you undo the last commit?", +10,493 votes, Community Wiki
In right bottom corner, number is reference to source (last slide)

Slide 5: Picture of bad SO question

"Forcing function to return if false" [locked, closed, off-topic], -154

Slide 6: Picture of Badges on SO

Picture of question badges

Slide 7: What is a question?

Questions can be generalized into either factoid or broad. Factoid questions usually only have a set amount of correct answers, whereas broad questions can have many answers which all are correct.

When using in education, questions are usually used as a learning tool to either help students learn something new, or through examination to evaluate your knowledge. For research, it could be the goal you are after, because you need to define a problem. You cannot just post a bunch of numbers, and say "These are my results". You need to ask the question "What are these results?", "What do these results tell me?", "What problems arise from these results?".

An interesting reverse situation is a game show from the early 90's called Jeopardy. The slogan for the show was "We have the answers, you have the questions". In this show, players were presented with an answer, and had to ask the question to which gave the answer. This could also be an interesting addition to learning, because what if the exam was not based on the teacher asking you a lot of questions? What if the exam was based on you asking the teacher questions to show that you understood the curriculum. Would you be able to ask hard enough questions to show that you grasped the curriculum?

The goal of Question classification (QC) is to categorize questions, since in most cases the goal is to find the answer to the question. By categorizing it, e.g. PERSON, LOCATION, DATE, you can reduce the amount of answers that could be related to the question.

You have WH-words, which are usually the first word in the question. Examples are "What, which, when, where, who, how, why". Some of which are harder to classify, because they are not as restrained as the others ("What", "Why", "How" and "Which").

N-grams is simply reducing the full text into fragments. These fragments can either be words or the characters in the word, and the 'N' represents the amount of fragments. One of the more known is the Bag-of-words (BOW), which is also called a unigram. BOW simply takes every word, counts the frequency and ignores the order. Bi-gram takes every second word, which means some order is kept. Meaning that for higher N-grams, the more focused you are on the sentence, the higher the N-gram should be.

If you want each word and sentence to be equal, you can remove the case-sensitivity. Because, when comparing characters, remember that large 'A' is not equal to the small 'a'. You can use semantics for word filtering, e.g. to remove duplicates or synonyms (word with same meaning). For instance WordNet, has a built in function for synonyms called synset. Hypernyms could also be of interest. Hypernyms are words that belong to the same category, sharing a parent-child relationship.

If you want a further word reduction, stemming is also an option. Stemming reduces the word to its base-form, e.g. crying would be converted into the word cry. Text can be split through tokenization, and you can also use stop words, which ignores words that are frequently used in the given language. You can also extract grammatical properties by using Part of Speech (POS).

Slide 8: Support Vector Machine (SVM)

SVM are good for solving regression and classification problems, where it attempts to solve a linear problem by using a hyperplane. It is mainly used for binary classification, where the two classes are separated by the hyperplane. Between the hyperplane and the classes, there are what is called a margin, which is the distance between the closest data points from the class and the hyperplane. These datapoints are called support vectors.

SVM consists of four kernels:

- Linear
- Radial Basis Function (RBF)
- Polynomial
- Sigmoid

Slide 9: Methodology

The data set that was used in this thesis was downloaded from the Stack Exchange Archive on March 30th. Originally, I used a data set from August 2015, but found that this was outdated, and I therefore downloaded a newer one. The data set contains XML files, where each element is a row containing the data from their database. Given that some of these files were several gigabytes in size, the only way to get access to the data was to load it into a MySQL database.

Since all the questions and answers are stored in the same table, these had to be separated, since it was time-consuming to run queries against it. Therefore two tables were created, having the same column as in posts, where one table contained all questions with score < 0, and the other questions with score > 0 (since questions with score of 0 was not of interest). A neat thing that does exist in python is Pandas. Pandas is a data analysis library, which includes something that is called DataFrame. The nice thing with DataFrame is that it not only contains the structure of the data as it is in the database, it is also supported by Scikit-learn as an array container.

For this development I used Python 3.5, and the development version of Scikit-learn. This was mainly because I had a lot of installation issues, and I had to install it from the GitHub repository. In addition, there was a lot of changes that were made in this version, so I decided to stick with it, instead of changing to the stable version.

Since all the questions were stored as HTML, they needed to be processed. This was done by creating a HTML parser class, and using BeautifulSoup to account for missing tags and "beautify" the HTML. When starting with this, I tried using the CountVectorizer. CountVectorizer takes each word in the passed text and counts the frequency of that word for the given text, and this is added to its vocabulary.

Since there was a lot of code examples, these had to be removed. This was done by checking for the <code> tag, and then simply removing the text. However, this meant that you lost the fact that this question contained code, so this was later replaced with the feature detector 'has_codeblock'. When also looking at the printed vocabulary, I saw that it contained a lot of hexadecimal and numerical values, so therefore these were also replaced with a feature detector called 'has_hexadecimal' and 'has_numerical'. Furthermore, there were a lot of dummy words used (an example is shown on the next slide), which had to be removed. The solution was the minimum document frequency

parameter, which can be used to ignore words that appear in either less than a given percentage of the documents or a given number of documents.

To be able to find out what could be good features for separating the good and the bad questions, 200 questions were looked at on SO (based on highest/lowest voted). As previously mentioned, it was easier to spot the bad questions (e.g. homework, bad code examples, etc).

For the homework words, I first attempted to use WordNet.synset, but it only gave three words: 'homework', 'prep', 'preparation'. I therefore looked at Thesaurus, which listed a lot more, and instead used my own defined dictionary. However, one issue was that the word 'assignment' could also be used in relation to programming terms (e.g. assignment operator), and I therefore had to split the Homework feature into two, one for homework words and one for the word 'assignment'.

I also saw that a lot of questions used links, either to link to external resource/tutorial or to show the "possible duplicate" flag.

SO also uses Tags for their questions, where a question can have from 1 to 5 tags total. Tags come in two variations, what I call external and attached. External are the list of tags (and their synonyms) that exists on the site, whereas attached are the ones the user has linked to their question. The problem is that the attached tags are wrapped as HTML, which had to be removed. Furthermore, they were all in lower-case, so all the questions needed to be converted to lower-case. In addition, they also needed to be sorted by word length, because there are two functionalities being used; Regular expression and string replacement. The regular expression looks for singular words only (e.g. 'C'), but since 'C++' counts the '+' as regex meaning 1 or more, string replacement had to be used. If the term 'C' was looked for, it would replace all single word entries only, meaning that occurrences of C++ would be 'has_tag'++. *also something about string replacement replacing occurrences inside the word for singular terms*

To avoid selecting some random parameter values, I used the default values from the Scikit-learn tutorials and GridSearchCV, which uses an exhaustive search. An exhaustive search means that it matches all the set values against each other to find the best match. As for the classifiers, two were selected, Support Vector Classification (SVC) and Stochastic Gradient Descent (SGD).

To have a replicable setting for the results, I used train_test_split from Scikit-learn, which splits the data into a set amount of training and test data. You can also assign a random state, and if the random state value is the same, then the results will be the same for each time you run it with the same data and settings. I also used cross-validation, which does its own validation on the training data, by splitting it into folds. The data is then trained on k-1 folds, and evaluated against the last fold.

Slide 10: Picture of dummy word

Title: Wordrap a very long string

Slide 11: Experiments and Results

A total of six different features were selected:

- Code samples
- Hexadecimal
- Homework (synonyms for homework and assignment)
- Links
- Numerical
- Tags (external and attached), external replaced to many "normal words"
- All features (using only homework and attached)
- 4 different experiments
 - Unprocessed data set vs. all singular feature, and all questions
 - Unprocessed data set vs. all singular features, and question occurrence only
 - Unprocessed data set vs. selected set of features only
 - Stochastic Gradient Descent (SGD) as classifier

Slide 12: Conclusion

- Stack overflow as a question quality metric
- Limitations and issues
- Further work
 - Code blocks, Links and Numerical as a feature set
 - Code analysis
 - Sentiment analysis
 - Version numbering

Slide 13: Demo

Show demo (if time)

Slide 14: Thanks for listening

Slide 15: References